

# HEAP: Reliable Assessment of BGP Hijacking Attacks

Johann Schlamp, Ralph Holz, Quentin Jacquemart, Georg Carle, and Ernst W. Biersack

**Abstract**—The detection of BGP prefix hijacking attacks has been the focus of research for more than a decade. However, the state-of-the-art techniques fall short of detecting more elaborate types of attack. To study such attacks, we devise a novel formalization of Internet routing, and apply this model to routing anomalies in order to establish a comprehensive attacker model. We use this model to precisely classify attacks and to evaluate their impact and detectability. We analyze the eligibility of attack tactics that suit an attacker's goals and demonstrate that related work mostly focuses on less impactful kinds of attacks. We further propose, implement, and test the *Hijacking Event Analysis Program (HEAP)*, a new approach to investigate hijacking alarms. Our approach is designed to seamlessly integrate with the previous work in order to reduce the high rates of false alarms inherent to these techniques. We leverage several unique data sources that can reliably disprove malicious intent. First, we make use of an Internet routing registry to derive business or organizational relationships between the parties involved in an event. Second, we use a topology-based reasoning algorithm to rule out events caused by legitimate operational practice. Finally, we use Internet-wide network scans to identify SSL/TLS-enabled hosts, which helps to identify non-malicious events by comparing public keys prior to and during an event. In our evaluation, we prove the effectiveness of our approach, and show that day-to-day routing anomalies are harmless for the most part. More importantly, we use HEAP to assess the validity of publicly reported alarms. We invite researchers to interface with HEAP in order to crosscheck and narrow down their hijacking alerts.

**Index Terms**—BGP hijacking, IRR analysis, SSL/TLS measurements, routing model.

## I. INTRODUCTION

THE Border Gateway Protocol (BGP) is today's standard for exchanging network routes between autonomous systems (ASes). Despite being vital to forward traffic on the Internet, BGP does not feature security mechanisms to validate route updates. Reports such as [1] and [2] have shown that attacks on BGP do occur and pose a real threat. Systems like

S-BGP [3] and RPKI [4] have been developed to add integrity protection and origin authentication to BGP. However, due to the considerable resources needed to deploy them, they are not widely used. Consequently, a number of mechanisms to detect routing attacks have been proposed [5]–[10].

Our contribution in this paper is twofold. First, we devise a novel formalization of Internet routing based on concepts from formal languages. With this model, attacks on BGP can be precisely formulated and classified according to their effect on the global routing table. This leads us to a comprehensive attacker model. We further discuss the motivation behind routing attacks and learn that common prefix hijacking offers no real benefit for an attacker apart from being destructive to the victim. More elaborate types of attacks aim to support sustained malicious activity with benefits for an attacker such as a chance to abuse networks to stage other attacks, send unsolicited email, or to impersonate a victim. We show that state-of-the-art detection techniques are not fully capable of dealing with the full spectrum of attacks. Surprisingly, most related work focuses on the less impactful attacks, and either fully neglects more effective and sustainable variants, or is of limited use due to a high rate of false positives.

In the second part of this paper, we present a scheme to assess the validity of generic hijacking alarms. This *Hijacking Event Analysis Program (HEAP)* is introduced as an automated system to reason about elaborate routing attacks. We leverage a carefully selected set of filters proposed in previous work [11], [12]. With HEAP, we strive to reliably identify *legitimate* events, i.e. hijacking alarms that are in fact false positives. To this end, we provide 1) administrative assurance obtained from Internet Routing Registries, 2) operational assurance based on insights into common routing practices, and 3) cryptographic assurance gained by comprehensive SSL/TLS measurements.

We evaluate HEAP for a set of common routing anomalies, so-called subMOAS conflicts, observed in BGP over the period of one month. Although a coarse approximation of real alarms, we learn that our system is highly effective in identifying legitimate events. We complement our findings with an instructive case study on routing anomalies for popular networks, which host the top one million web sites. Attackers have much interest to launch malicious activities from such networks as they can profit from their good reputation. We show that HEAP naturally benefits from this circumstance and yields very good results: We are able to rule out attacks for more than 80% of corresponding routing anomalies.

By studying day-to-day anomalies without intrinsic evidence for an attack, we establish a base line for our

Manuscript received September 11, 2015; revised January 20, 2016; accepted February 15, 2016. Date of publication April 27, 2016; date of current version June 6, 2016. This work was supported by UCN@Sophia Labex.

J. Schlamp and G. Carle are with the Technical University of Munich, Garching 85748, Germany (e-mail: schlamp@net.in.tum.de; carle@net.in.tum.de).

R. Holz was with Data61/CSIRO, Eveleigh, NSW 2015, Australia. He is now with The University of Sydney, Sydney, NSW 2006, Australia (e-mail: ralph.holz@sydney.edu.au).

Q. Jacquemart was with Eurécom, Sophia Antipolis 06904, France. He is now with the University of Nice Sophia Antipolis, Nice 06000, France (e-mail: quentin.jacquemart@unice.fr).

E. W. Biersack was with Eurécom, Sophia Antipolis, France. He is now with Caipry, Valbonne 06560, France (e-mail: erbi@e-biersack.eu).

Color versions of one or more of the figures in this paper are available online at <http://ieeexplore.ieee.org>.

Digital Object Identifier 10.1109/JSAC.2016.2558978

legitimization capabilities. To demonstrate the effectiveness of HEAP, we further apply our assessment scheme to publicly reported hijacking alarms and learn that even such a set of highly suspicious events still contains nearly 10% of false positives. We find our results highly encouraging and conclude that HEAP is suitable to assess a variety of hijacking alarms, including those of related work that inherently exhibit high rates of false positives. Hence, we call on researchers to continuously feed our system with their conjectural alarms.

The remainder of this paper is organised as follows. Section II presents our formalized attacker model. We discuss and assess related work in Section III. Our methodology for HEAP is presented in Section IV, followed by an evaluation of the approach in Section V.

## II. A COMPREHENSIVE ATTACKER MODEL

We present a novel formalization of Internet routing based on concepts of formal languages. With this model, routing anomalies can be precisely expressed, classified, and further assessed with respect to impact and detectability. Previous work mostly relies on informal and often inconsistent definitions. We believe that our routing model can improve on this situation. It may serve as a well-founded basis for a broad spectrum of future analyses on Internet routing.

### A. Formalization of Internet Routing

As per definition, a formal language is a set of strings of symbols constrained by specific rules. Analogous to that, Internet routing can be represented as the set  $\mathcal{L}$  of all active BGP routes in the global routing system, i.e. the set of AS paths from all vantage points towards all advertised IP prefixes. In this model, a routing attack is then defined by an attacker extending  $\mathcal{L}$  by forged routes.

1) *Preliminaries:* Let  $\Sigma_{AS}$  be the set of all ASes,  $\Pi$  the set of all IP addresses,  $p \subset \Pi$  an IP prefix, and  $p' \subset p$  a more specific prefix of  $p$ . Let further be  $(w, p) \in \Sigma_{AS}^* \times \Pi$  a route with an AS path  $w \in \Sigma_{AS}^*$ , i.e. an arbitrary concatenation of ASes, to a prefix  $p \in \Pi$ , in the following denoted  $r = wp$ . Then, we define  $\mathcal{L} \subset \Sigma_{AS}^* \times \Pi$  as the set of *active* routes to all advertised prefixes in the global routing system, i.e. the set of all observable routes.  $\mathcal{L}(p) \subset \mathcal{L}$  denotes the subset of routes to a given prefix  $p \subset \Pi$ , such that

$$\mathcal{L}(p) = \{wup \in \mathcal{L} \mid w \in \Sigma_{AS}^* ; u \in \Sigma_{AS} ; o \in \Sigma_{AS}\}$$

with  $w$  being an AS subpath and  $u$  the upstream AS of the origin AS  $o$ . For a given route  $r$  and a subprefix  $p' \subset p$ , we postulate  $r \in \mathcal{L}(p) \Rightarrow r \in \mathcal{L}(p')$  as a corollary, since routes to less specific prefixes also cover more specific prefixes. Note that the converse is false. Further,  $\mathcal{L}_P \subset \mathcal{L}$  denotes the set of all observable routes from a set of observation points  $P \subset \Sigma_{AS}$ .  $\Pi_o$  denotes the set of IP addresses advertised by an AS  $o \in \Sigma_{AS}$ , i.e. the union of its advertised prefixes. Then, the set of origin ASes for a prefix  $p \subset \Pi$  is given by

$$O(p) = \{o \in \Sigma_{AS} \mid p \subset \Pi_o\}.$$

Consistently, the set of origin ASes  $O(p')$  for a subprefix  $p' \subset p$  comprises the origin ASes for less specific prefixes

such that

$$O(p') = \{o \in \Sigma_{AS} \mid p' \subset \Pi_o\} = \bigcup_{p \supseteq p'} O(p),$$

since these ASes effectively originate routes to the particular network  $p'$ . Note again that the converse is false. The set of upstream AS neighbors for an AS  $o \in \Sigma_{AS}$  is given by

$$U(o) = \{u \in \Sigma_{AS} \mid wuop \in \mathcal{L} \text{ such that } w \in \Sigma_{AS}^* ; p \subset \Pi_o\}.$$

Due to *best path selection* in BGP, the number of routes from an observation point  $s \in \Sigma_{AS}$  to a particular prefix  $p \subset \Pi$  is generally bound by the number of neighboring ASes of  $s$ , i.e.  $|\mathcal{L}_s(p)| \leq |U(s)|$  holds. In the following, we reuse the unary operator  $|\cdot|$  to indicate the number of routes in a set  $\mathcal{O} \subseteq \mathcal{L}$ , the length of a route  $r \in \mathcal{L}$  or a subpath  $w \in \Sigma_{AS}^*$ , and the number of ASes in a set  $S \subseteq \Sigma_{AS}$ .

2) *Definition of Routing Attacks:* We define routing attacks as an attacker extending the global set of BGP routes  $\mathcal{L}$  by forged routes  $\mathcal{F}$ . Their purpose is to manipulate existing routes in order to re-route traffic flows or to take over a victim's Internet resources. In general, such incidents are called *hijacking attacks*. Typically, these attacks lead to topological changes in the Internet, which can be observed by neutral BGP speakers.

Our attacker is assumed to be capable of injecting arbitrary BGP messages into the global routing system, i.e. he operates a BGP router and maintains a BGP session to at least one upstream provider. We assume that the attacker is not hindered by local filters or other validation mechanisms employed by his upstream provider. Instead, the upstream AS indifferently redistributes all update messages to its peers, which thus may propagate throughout the Internet. An observation point shall be in place to monitor the propagation of BGP messages. It is worth mentioning that data packets do not necessarily traverse all ASes in a given path, since an attacker may craft BGP messages with a forged AS path. Further, route updates with less attractive paths may not reach a particular observation point due to best path selection in BGP. Without loss of generality, an omnipresent observation point to observe the set of all active routes  $\mathcal{L}$  is assumed for the following definitions.

In the following, we denote an attacker's AS  $a \in \Sigma_{AS}$  and his victim's AS  $v \in \Sigma_{AS}$ . Further, a victim's prefix is given by  $p_v \subset \Pi_v$ . Then, a generic routing attack on  $p_v$  is defined by an attacker injecting forged routes  $\mathcal{F}_a$  into the routing system, such that the altered set of globally visible routes  $\hat{\mathcal{L}}(p'_v)$  is given by

$$\hat{\mathcal{L}}(p'_v) = \mathcal{L}(p_v) \cup \mathcal{F}_a(p'_v) \text{ with } p'_v \subseteq p_v.$$

3) *Impact Analysis:* In BGP, the impact of a hijacking attack generally depends on a best path selection process. In particular, shorter AS paths are preferred over longer ones, although policy-induced exceptions on a case-by-case basis may exist. With respect to packet forwarding, routes to longer IP prefixes prevail. Assuming the ambition to forge globally accepted routes, an attacker thus succeeds if his routes towards a victim's network are considered best by a vast majority of




Fig. 1. Origin relocation attacks in BGP.

Internet participants. In practice, an attacker needs to ensure that his bogus routes  $\mathcal{F}_a(p'_v)$  are either

- 1) unrivalled in terms of competitive routes, i.e.  $|\mathcal{L}(p'_v)| = 0$ ,
- 2) shortest from a global perspective, i.e.  
 $\forall r \in \mathcal{L}(p'_v), r_a \in \mathcal{F}_a(p'_v) : |r_a| < |r|$ , or
- 3) more specific than all others, i.e.  
 $\forall p''_v \subseteq p'_v \subset p_v : \mathcal{L}(p''_v) = \mathcal{L}(p'_v) = \mathcal{L}(p_v)$ .

As a consequence, the prospects of identifying an attack naturally depend on the significance of topological changes in  $\hat{\mathcal{L}}$ , i.e. on abnormal changes to the sets of origin ASes  $\hat{O}$  and upstream ASes  $\hat{U}$  for a victim's network.

### B. Classification of Attacks

BGP-based attacks aim to inject falsified protocol messages into the global routing system, which may lead to topological changes in the Internet. Depending on the characteristics of such changes, hijacking attacks can be classified into several subtypes with differing tactical value.

1) *Examples*: In addition to the formalized model, the topology in Figure 1 is used to exemplify different types of attacks. *Mallory* thereby denotes an attacker, and her autonomous system respectively. *Oscar* and *Paul* are *Mallory's* upstream providers who do not validate BGP messages. The observation point *OP* receives route updates that propagate through the Internet, which are herein after simplified to the topology-relevant attributes of BGP messages, namely the IP prefix, also called NLRI, and the AS path, referred to as AS\_PATH. The following expression illustrates such an update observed at *OP*:

$$OP: AS_*^K \leftarrow Carol \leftarrow Alice \ll P \text{ legitimate}$$

Originating ( $\ll$ ) at *Alice*, a route update for the prefix *P* traverses *Carol* and a series of *K* ASes  $AS_*^K$  to reach the observation point *OP*. For the sake of clarity, temporary convergence effects within BGP are ignored. *Alice* serves as a victim for different kinds of attacks, *Carol* and *Dave* provide upstream connectivity for *Alice*.

2) *Hijacking Attacks in Practice*: Given a standard router with an already established point-to-point IP connection to the router of an upstream provider, it is surprisingly easy to participate in BGP and to originate an IP prefix. Consider the following BGP configuration of *Alice's* router named *rtr*, which actually represents a minimal working example:

```
rtr(config)# router bgp alice
rtr(config-router)# neighbor X.X.X.X remote-as carol
rtr(config-router)# network 10.1.0.0/16
```

In this example, *Alice* opens a BGP session to her neighbor router X.X.X.X, which is operated by *Carol*, and advertises direct reachability of 10.1.0.0/16. This information will be redistributed by *Carol* to her neighbors, and may eventually propagate to all BGP routers connected to the Internet. It is as easy for an attacker to originate arbitrary IP prefixes, irrespectively of whether legitimate routes to these IP prefixes already exist. Depending on the specifics of the attacker's approach, the original routes will be partially or entirely overridden in the global routing system. Hence, this type of attack results in an *origin relocation* of a victim's network.

3) *Prefix Hijacking*: The most basic form of origin relocation attacks is *prefix hijacking*. An attacker thereby originates a victim's prefix at his own AS, in principle in the same way as illustrated above. The resulting forged routes compete with the victim's concurrent announcements. The following definition formulates this attack scenario:

$$\hat{\mathcal{L}}(p_v) = \{wuvp_v \mid w \in \Sigma_{AS}^*; u \in U(v)\} \cup \text{legitimate} \\ \{wuap_v \mid w \in \Sigma_{AS}^*; u \in U(a)\} \quad \text{forged}$$

$$\text{with } \hat{O}(p_v) = O(p_v) \cup \{a\}$$

$$\text{and } \hat{U}(v) = U(v)$$

The set of origin ASes  $\hat{O}(p_v)$  for the prefix  $p_v$  now comprises two ASes, while the set of the victim's upstream ASes  $\hat{U}(v)$  remains unchanged. In literature, a situation with  $|O(p)| > 1$  is often called a *multi-origin AS (MOAS)*. Given the exemplary topology in Figure 1, *Mallory* may craft a BGP update message as listed below. At the same time, legitimate paths advertised by *Alice* are present in the global routing table.

*Example 1: Prefix Hijacking.*

$$OP: AS_*^K \leftarrow Carol \leftarrow Alice \ll 10.1.0.0/16 \text{ legitimate}$$

$$OP: AS_*^L \leftarrow Dave \leftarrow Alice \ll 10.1.0.0/16 \text{ legitimate}$$

$$OP: AS_*^M \leftarrow Oscar \leftarrow \boxed{\text{Mallory}} \ll 10.1.0.0/16 \text{ forged}$$

As shorter AS paths are generally preferred over longer ones, the attack is likely to succeed for observation points *s* where  $M < \min(K, L)$  holds, and for

$$\{s \in \Sigma_{AS} \mid \forall w_v p_v \in \mathcal{L}_s(p_v) \exists w_a p_v \in \mathcal{F}_a(p_v) : |w_a| < |w_v|\}$$

respectively. However, it is safe to assume that clients that are topologically close to *Carol* or *Dave* still reach the victim *Alice*, since shorter legitimate routes take precedence. The Internet thus decomposes into two disjoint parts: one part that is affected by the forged announcement, in literature often called the *poisoned* part, and one that remains unaffected.

4) *Subprefix Hijacking*: To overcome the limited impact inherent to prefix hijacking, an attacker can leverage longest prefix matching in IP routing with so-called *subprefix hijacking*. To this end, the attacker originates a subprefix  $p'_v \subset p_v$  at his AS, thereby injecting a new set of routes  $\mathcal{F}_a(p'_v)$  into the global routing system as given by:

$$\hat{\mathcal{L}}(p_v) = \{wuvp_v \mid w \in \Sigma_{AS}^*; u \in U(v)\} \cup \text{legitimate} \\ \{wuap'_v \mid w \in \Sigma_{AS}^*; u \in U(a)\} \quad \text{forged}$$

$$\text{with } \hat{O}(p_v) = O(p_v), \hat{O}(p'_v) = O(p_v) \cup \{a\}$$

$$\text{and } \hat{U}(v) = U(v)$$

Subprefix hijacking attacks generally have global impact, since routes to the more specific prefix  $p' \subset p$  dominate. Such incidents with  $|O(p)| > 0$  and  $|O(p') \setminus O(p)| > 0$  are called *subprefix multi-origin AS* (*subMOAS*).

Note that the victim might readily advertise routes to a prefix and a corresponding subprefix concurrently to the attacker's subprefix route, i.e. condition 3) in Subsection II-A3 does not hold since  $\mathcal{L}(p_v) \subset \mathcal{L}(p'_v)$ . In this case, the event can also be considered a MOAS event. Otherwise, it is called a *strict subMOAS*, and subprefix hijacking respectively. We assume this variant in the following. While virtually all Internet participants are affected by strict subprefix hijacking, it may be tempting to conclude that only part of the victim's network, i.e. subprefixes, can be taken over. As a matter of fact, this is not the case. An attacker can easily craft multiple update messages such that the set of forged routes  $\mathcal{F}_a(p_v)$  fully covers the prefix  $p_v$  with more specific routes:

$$\mathcal{F}_a(p_v) = \bigcup_{p'_v \in \Pi} \mathcal{F}_a(p'_v) \text{ such that } p_v = \bigcup p'_v.$$

With respect to Figure 1, *Mallory* could thus inject the following BGP routes:

*Example 2: Subprefix Hijacking.*

$OP : AS_*^K \leftarrow Carol \leftarrow Alice \ll 10.1.0.0/16$  *legitimate*

$OP : AS_*^L \leftarrow Dave \leftarrow Alice \ll 10.1.0.0/16$  *legitimate*

$OP : AS_*^M \leftarrow Oscar \leftarrow \boxed{Mallory} \ll 10.1.0.0/17$  **forged**

$OP : AS_*^M \leftarrow Oscar \leftarrow \boxed{Mallory} \ll 10.1.128.0/17$  **forged**

By advertising a victim's network with BGP updates split up into multiple longer prefixes, as given by  $10.1.0.0/17$  and  $10.1.128.0/17$  in the example above, subprefix hijacking can be as extensive as regular prefix hijacking, with the advantage of globally preferred routes at the same time. Notwithstanding this possibility, an attacker might be satisfied with hijacking individual subnets of high value only.

5) *Other Types of Attack:* Origin relocation attacks as discussed above lead to noticeable changes in the set of origins for a victim's prefix. In contrast, *route diversion* attacks aim at the manipulation of AS paths towards a victim. In its basic form, an attacker hijacks a victim's prefixes and its AS, which effectively disguises the attack by hiding the attacker's own AS. Moreover, tailored attacks can be derived to impersonate a victim from an administrative point of view by stealthily hijacking abandoned Internet resources. We have studied such *hidden takeover* attacks in great detail in previous work [2], [13]. An even more sophisticated attack based on AS path manipulation aims at stealthily intercepting a victim's traffic while maintaining the victim's connectivity. Such *man-in-the-middle* attacks require a stable backhaul link from the attacker to the victim to forward eavesdropped packets. Interestingly, BGP itself, and its loop prevention mechanism respectively, can be leveraged to accomplish this task. Our routing model allows to formalize and study these types of attacks. For this paper, however, we consider them beyond scope.

### C. Motivation Behind Hijacking Attacks

Hijacking attacks may serve a variety of purposes, which can be divided into *destructive* and *abusive* variants.

An obvious intent is to inflict damage to a victim's operations by disrupting network connectivity, which is often called *blackholing* in literature. Furthermore, hijacked networks can be abused for malicious short-term activities, for instance to launch a fast, temporary, and massive spam campaign. Abusive actions may aim at hosting illegal services, like phishing web sites, or serve to establish a stable base of operations for subsequent attacks, e.g. to command botnets from a safe-house network. More sophisticated attacks aim at compromising a victim's reputation by carrying out illicit actions. Lastly, attacks can be tailored to break confidentiality or integrity of a particular victim's communications by means of interception.

Prefix hijacking attacks *partially* disrupt a victim's connectivity, but are of limited use in other respects, like hosting malicious services, since a significant part of the Internet might still prefer the victim's routes. In contrast, subprefix hijacking attacks are capable of breaking communications entirely, i.e. all hosts inside the victim's network become globally unreachable. Hence, this type of attack is also useful for launching illegal operations from a hijacked network, and an important element for more sophisticated attacks like AS hijacking or man-in-the-middle interception. Having no real-world benefit apart from partially disconnecting a victim's hosts from the Internet, it is surprising to see that the state-of-the-art focuses primarily on prefix hijacking.

## III. RELATED WORK

There is a large body of literature on the detection of BGP-based routing attacks. Corresponding techniques can be divided into control-plane and data-plane techniques, with hybrid approaches emerging recently.

### A. State-of-the-Art Detection

Possibly the first attempt to detect hijacking attacks was presented with PHAS [5]. It is a control-plane technique focusing exclusively on reporting MOAS conflicts. In an effort to reduce high rates of false positives, PHAS utilizes an adaptive time-window that prevents recurring alerts. The authors of [6] provide further heuristics to assess MOAS conflicts with respect to compliance with economy-based routing policies. Careful tuning of heuristic parameters is necessary to yield suitable detection results. The approach thus tends to reduce false alarms at the cost of an increased rate of false negatives. LOCK [14] offers to pinpoint attackers in the AS-level topology. The Buddyguard system [15] uses a learning-based approach to detect abnormal routing changes. The creators of BGPmon.net [16] provide MOAS and subMOAS alarms in real-time, but do not publish details about their methodology.

In [7], a light-weight distributed measurement scheme (LWDS) is proposed. This data-plane technique is based on the assumption that path measurements yield stable path lengths for a majority of networks. Major violations of this conjecture hint at a suspicious change in network location if a set of reference points close to this network remains unaffected. Due to its dependency on suitable vantage points and reference nodes, it is difficult to deploy the scheme on a larger scale. A comparable technique

TABLE I  
COMPARISON OF POPULAR STATE-OF-THE-ART DETECTION SYSTEMS

	Prefix Hijacking	Subprefix Hijacking	Applicability	Public Interface
<i>control-plane detection</i>				
PHAS [5]	+	o	o ( <i>operator-level</i> )	no
Bogus Routes [6]	+	o	+	no
BGPmon.net [16]	x	x	+	yes
<i>data-plane detection</i>				
LWDS [7]	+	+	– ( <i>incident-level</i> )	no
iSpy [8]	+	–	o ( <i>operator-level</i> )	no
<i>hybrid techniques</i>				
Fingerprints [9]	+	o	+	no
Argus [10]	+	o	+	yes
<b>HEAP (our system)</b>	<b>o</b>	<b>++</b>	<b>+</b>	<b>yes</b>

++/+: practical usefulness    o: theoretic applicability  
–: not supported    x: no details published

that leverages latency measurements from multiple vantage points is proposed in [17]. The StrobeLight system [18] detects hijacking attacks in a similar way, though on a per-operator basis only. A more popular operator-centric approach is given with iSpy [8]. By carrying out data-plane measurements from an operator’s network to major transit ASes, attacks are detected based on specific outage patterns that reflect the partitioning of the Internet into affected and unaffected parts (refer to Subsection-II-B3). iSpy cannot differentiate between subprefix hijacking and temporary link failures or congestion near the victim’s network.

One of the first hijacking detection systems to include both passive and active measurements is a fingerprint-based approach [9]. This system utilizes a monitoring scheme for BGP combined with active measurements to identify two disjoint parts of the Internet as outlined above. For subprefix hijacking, where no such partitioning occurs (see Subsection II-B4), several heuristics are proposed. A similar approach is realized with the Argus system [10]. This framework detects anomalies in BGP and confirms hijacking attacks by carrying out measurements from public looking glasses to identify a poisoned part of the Internet. The authors of [19] extend the set of active scans, while [20] proposes additional types of fingerprints. A complementary approach to assess immediate effects of control-plane anomalies onto the data-plane is to perform measurements during and after an event [21].

### B. Assessment and Comparison

Looking at Table I, it is surprising to see that related work concentrates its effort on the detection of prefix hijacking attacks, despite the fact that this kind of attack is of very limited use for an attacker in practice. As a matter of fact, subprefix hijacking attacks offer more useful means for a broader area of operations, but are mostly neglected in state-of-the-art techniques. HEAP, our own contribution, aims to improve on this situation by providing reliable means to assess alarms relating to this specific type of attack.

## IV. HEAP: A REAL-TIME FRAMEWORK

We propose a novel *Hijacking Event Analysis Program (HEAP)* to reliably assess hijacking alarms.

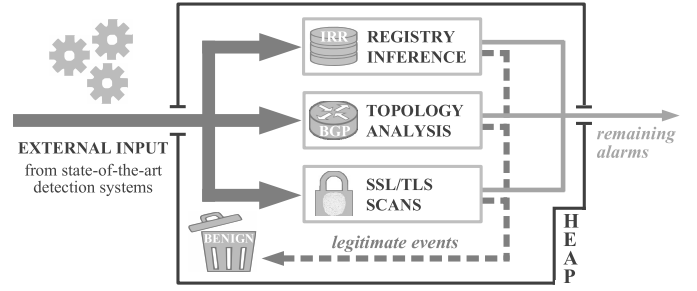


Fig. 2. The Hijacking Event Analysis Program (HEAP).

We combine techniques from prior work [11], [12] to identify incidents with legitimate cause. Our goal is *not* to develop a new alarm system. Instead, HEAP is designed to receive input from available detection systems to reduce their rate of false alarms.

### A. System Architecture

HEAP leverages three distinct data sources to assess hijacking events. Our main assumption here is that an attacker is capable to hijack networks in BGP, but cannot alter orthogonal data sources that relate to the operation of hijacked networks. Hence, we are able to rule out an attack if these data sources legitimize a suspicious routing anomaly. First, Internet Routing Registries (IRR) are utilized to infer legitimate business relationships between an attacker and his alleged victim. Second, a topological analysis is carried out in order to identify benign anomalies resulting from common operational practices. And lastly, SSL/TLS measurements yield cryptographic assurance that traffic to a supposedly hijacked network is still routed to the alleged victim.

Figure 2 illustrates the main workflow within HEAP. Given external alarms fed into the system, legitimate events are identified and eliminated as false positives based on the aforementioned data sources. Note that the SSL/TLS component needs a tight coupling to external systems that provide us with alarms since corresponding scans have to be carried out in response to the input received. The remaining events are highly suspicious indications for an attack, even if we take into account that we cannot make further assumptions about their nature. This is for two reasons: 1) the input source already provides potential hijacking candidates, and 2) none of our filter techniques yields evidence for a legitimate cause. In the following, we will show that this is indeed improbable for benign events. The remaining alarms lend themselves well to manual inspection, with a rich set of background information readily available from the individual analysis steps.

### B. Filtering Methodology

All filters applied by HEAP are executed concurrently. HEAP is easily extensible, i.e. additional filters can be incorporated without difficulty. Three independent techniques to eliminate legitimate alarms have been implemented so far.

1) *Utilizing IRR Databases:* Regional Internet Registrars (RIRs) maintain so-called Internet Routing Registries (IRR), i.e. databases that contain information

TABLE II  
DATA STORED IN OUR GRAPH DATABASE. AUGUST, 2015

instance	AfriNIC		APNIC		<sup>1</sup> ARIN		LACNIC		RIPE	
	nodes	relations	nodes	relations	nodes	relations	nodes	relations	nodes	relations
MNTNER ← <i>maintained_by</i> [ * ]	2,624	—	20,129	—	n/a	—	n/a	—	53,670	—
	—	133,186	—	1,919,397	—	n/a	—	n/a	—	5,620,385
ORGANISATION ← <i>org</i> [ * ]	1,877	—	n/a	—	2,976,707	—	n/a	—	90,102	—
	—	32,476	—	n/a	—	3,536,502	—	n/a	—	249,319
AUT-NUM ← <i>origin</i> ROUTE ← <i>import</i> AUT-NUM	1,239	—	9,485	—	24,939	—	5,193	—	29,206	—
	—	464	—	216,865	—	<sup>2</sup> 583,296	—	n/a	—	279,532
	—	6	—	10,734	—	n/a	—	n/a	—	228,509
INETNUM	85,672	—	924,584	—	2,910,623	—	342,104	—	3,995,522	—
ROUTE	443	—	97,858	—	<sup>2</sup> 600,940	—	n/a	—	267,216	—

<sup>1</sup> ARIN's object identifiers can be directly mapped to RIPE's schema (e.g. ASHandle → AUT-NUM).

<sup>2</sup> Implicitly given in ARIN's INETNUM objects (via OriginAS attributes).

pertaining to the management of Internet resources. A recent study [22] matched prefixes and ASes observed in BGP and IRRs by looking for appropriate database objects. We provide a more generalized set of inference rules to identify benign routing events, which take into account multiple prefix origins observed in BGP as well as complex relationships between affected prefixes and suspicious ASes. The fundamental assumption behind our approach is that an attacker does not have the credentials to change an IRR database in order to cover his attack. To disprove an attack, we accordingly look for legitimizing database relations between the entities involved in a hijacking alarm, e.g. for a common organisation referenced by two ASes. To this end, we download and evaluate snapshots of the IRR databases, which are provided by RIRs on a daily basis. We use a graph database to store the extracted information using the schema presented in Table II and track all changes over time. Note that IRR databases are updated by individual resource holders and can thus be outdated or even hold conflicting information. Our filter accounts for this by strictly searching for legitimizing relationships without drawing any conclusions in their absence.

To assess a given hijacking alarm, we map the affected AS numbers and prefixes to resource objects, i.e. graph nodes, in our graph database. We then traverse the graph along a path of legitimizing relations that document a right to use, which are given by AUT-NUM and INETNUM objects linked by *import*, *origin*, *maintained\_by* or *org* relations. We look for such paths between a) two affected ASes, or b) a prefix and its origin AS. If we succeed with a), we can infer a valid business relationship between the victim and the suspected attacker. If we succeed with b), the suspected attacker holds ownership rights for the prefix and is thus authorized to originate the prefix from his AS. Compared to our previous work [11], we extended this filter to support the IRR databases provided by all five RIRs. Note that we transform all IRR databases into the RIPE data model as presented in Figure 3, since it is most consistent and represents a superset of all available information. In fact, AfriNIC and APNIC already store their data in a similar format and can thus be directly processed by our filter. LACNIC and ARIN utilize their own data models, which can be converted in spite of some missing data points.

RIPE-based IRR databases model access rights with the help of MNTNER objects. Only maintainers with valid credentials

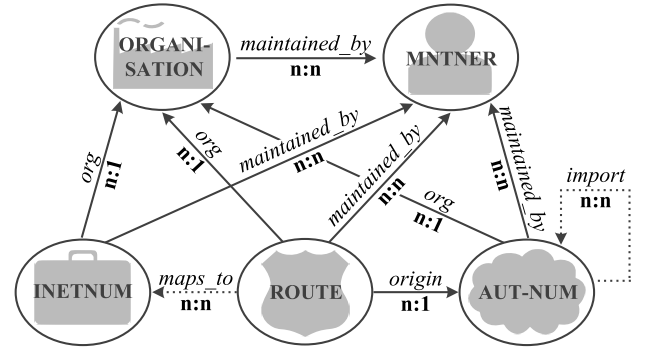


Fig. 3. Relevant data objects and relations in IRR databases.

can modify or delete objects. For any object, this is expressed by adding a *maintained\_by* reference pointing to the respective MNTNER object. ORGANISATION objects are mainly used to provide administrative contact details. For privacy reasons, most IRR database snapshots do not include details, but unique references to these objects are preserved. INETNUM objects document allocated or assigned IPv4 prefixes managed by the respective RIR. AUT-NUM objects represent AS numbers and may be referenced as the *origin* of ROUTE objects. Such ROUTE objects are created by resource holders and are used to document or confirm intended prefix announcements by specific ASes. To create a ROUTE object, the resource holder needs to provide valid credentials for the respective INETNUM and AUT-NUM objects. A corresponding *maps\_to* relation is derived by our parsing algorithm, as is the case with *import* relations deduced from free-text description fields, which are often used to model routing policies in the so-called Routing Policy Specification Language (RPSL). When resources are deleted from a database, RPSL definitions may still reference (now) non-existing ASes. We account for this by tracking such orphaned *import* relations.

Table II provides further details on selected objects that are relevant to our approach. Our combined database holds more than 15 million nodes and 45 million relations extracted from the five individual IRR databases. Entries marked with *n/a* are not available in the respective database snapshots. For the RIPE database, for instance, we can see that less than 55,000 MNTNER objects share more than 5 million incoming *maintained\_by* references. Although optional, roughly 90,000 ORGANISATION objects are referenced by 250,000 other

objects. About 280,000 ROUTE objects bind prefix announcements to less than 30,000 AUT-NUM objects. Furthermore, these AUT-NUM objects document nearly 230,000 *import* routing policies. We will see that this rich information allows our filter to be highly effective—except in the case of LACNIC, where none of the necessary cross-referencing objects are provided in the daily database snapshots.

2) *Topology Reasoning*: The next filter to legitimize routing anomalies is a topology-based reasoning algorithm. The key idea is that an attacker is unlikely to hijack his own upstream provider. This assumption is based on the fact that the attacker’s malicious BGP updates need to propagate via this upstream provider, who could simply counter an attack by filtering them out. As a consequence, we can rule out an attack if the suspected attacker resides in the downstream AS path of his victim.

To identify such benign anomalies, we utilize BGP collectors to extract all AS paths that lead to the affected prefixes and ASes respectively. If we do not find any AS path that contains both the attacker’s and the victim’s AS, we cannot draw any further conclusions. The same is true for AS paths in which the attacker is located upstream of his victim. In contrast, we can infer a legitimate cause of the anomaly if the attacker is actually located downstream of the victim, i.e. if we find a particular AS path in which the victim’s AS precedes the attacker’s AS. In this case, we can rule out malicious intent.

Such benign situations might occur, for instance, if smaller organizations obtain Internet connectivity and an IP prefix from a larger carrier. Other reasons can be static routes invisible to BGP, imperfect multihoming setups, or even misconfiguration. In our evaluation, we will see that a significant part of day-to-day routing anomalies is caused by such topological constellations.

3) *Cryptographic Assurance With SSL/TLS*: We use a final strong filter that is based on our regular Internet-wide scans of SSL/TLS protocols (refer to [23] for further details). For any given hijacking alarm concerning a certain IP prefix, we verify if affected SSL/TLS hosts present the same public key before and during the event. We make the assumption that an attacker cannot gain access to the private keys of a victim’s hosts, and thus cannot perform successful SSL/TLS handshakes. We conclude that such cases cannot be attacks.

A prerequisite for this filter is a *ground truth scan* to obtain a known-correct mapping from IP addresses to public keys that are used on corresponding machines. Given such a ground truth data set, we can carry out *validation scans* to hosts in a prefix that relates to a hijacking alarm and compare the retrieved public keys. Note that it is imperative for these scans to be executed in a timely manner, i.e. we need to compare public keys during the life time of an event. A tight coupling to the alarming system is dispensable if we can retroactively ascertain that an event lasted for the entire duration of a corresponding scan. Since our system is designed to assess subprefix hijacking attacks, which affect the Internet as a whole (refer to Subsection II-A3), the vantage point for our SSL/TLS measurements can be chosen freely. For this paper, we employed a scanning machine hosted at our university in Munich (AS56357).

TABLE III  
SCANNED SSL/TLS HOSTS FOR OUR GROUND TRUTH DATA SET. ALL MEASUREMENTS WERE CARRIED OUT IN JULY, 2015

	port	time	port open	handshake	in %
<b>Implicit SSL/TLS</b>	27d		<b>72,546,563</b>	<b>38,146,816</b>	<b>52.58%</b>
HTTPS	443	10d	42,676,912	27,252,853	63.85%
SMTPS	465	2d	7,234,817	3,437,382	47.51%
IMAPS	993	3d	6,297,805	4,121,108	65.43%
POP3S	995	3d	5,186,724	2,797,300	53.93%
FTPS	990	2d	2,657,680	344,400	12.95%
LDAPS	636	2d	2,273,771	112,978	4.96%
XMPPS/CLIENT	5223	2d	2,223,994	70,441	3.16%
XMPPS/SERVER	5270	1d	2,046,204	1,693	0.08%
IRCS	6697	2d	1,948,656	8,661	0.44%
<b>Explicit SSL/TLS</b>	9d		<b>51,768,705</b>	<b>18,316,920</b>	<b>35.38%</b>
FTP/STARTTLS	21	2d	14,493,966	2,939,048	20.27%
SMTP/STARTTLS	25	2d	12,488,000	3,848,843	30.82%
POP3/STARTTLS	110	1d	8,930,688	4,074,211	45.62%
IMAP/STARTTLS	143	2d	8,006,617	4,076,809	50.91%
SUBMISSION/STARTTLS	587	2d	7,849,434	3,378,009	43.03%
<b>Total SSL/TLS scans</b>	36d		<b>124,315,268</b>	<b>56,463,736</b>	<b>45.42%</b>

Compared to our previous work [11], we greatly extended the ground truth scans to a variety of popular SSL/TLS-based protocols. Table III shows all scans that were carried out for this work. In many cases, we scanned for both TLS and STARTTLS, a common extension to network protocols that allows for opportunistic use of TLS. It is instructive to see that the use of TLS varies greatly between application-layer protocols. An open, dedicated port does not imply support for TLS per se. Due to the marginal contributions that our scans of XMPPS and IRCs provided, we have since stopped scanning these protocols.

On total, we tried to open connections to 124,315,268 individual ports. For successful SSL/TLS handshakes, we downloaded the certificate and extracted the public key. Note that we only consider keys that were unique across the whole dataset. This precaution eliminates the risk of falsely legitimizing events imposed by default certificates. Such certificates often ship with popular web server software or with SSL/TLS-enabled devices, and could thus be presented by an attacker as well. We relax this condition only where the same key is presented by a single host for multiple protocols. This finally yields a total of 12,800,474 available keys, which were presented by a total of 8,402,023 different hosts.

### C. Applicability

Our approach works best for the assessment of subprefix hijacking alarms. Attacks that build upon the manipulation of AS paths, like AS hijacking, for instance, can be assessed with HEAP as well. Due to a general lack of initially suspicious input events, however, we exclude this kind of attack from our analysis.

Ordinary prefix hijacking attacks, and MOAS conflicts respectively, impose limitations to our SSL/TLS filter, since we cannot assure that measurements reach a supposedly hijacked network. With the Internet decomposing into two disjoint parts as discussed in Subsection II-B3, our SSL/TLS scans might



reach either part, which prevents reliable conclusions. Nevertheless, we can deactivate the SSL/TLS filter for such cases. BGP-based man-in-the-middle attacks [24] are especially hard to identify [25]. In an interception scenario, in which an attacker is able to forward our active scans to the victim, the SSL/TLS filter would wrongly legitimize the incident. Hence, these attacks are left for future studies.

We acknowledge that our approach depends on external input, thus it is arguably not a full-fledged detection system. In our evaluation, however, we show that we arrive at remarkable validation results even for a superset of potential alarms.

## V. EVALUATION

Most of the proposed detection techniques in previous work (see Table I) do not offer publicly available interfaces yet. We compensate for the resulting lack of real alarms by studying common subMOAS conflicts observed in BGP. Such cases occur numerous times per day and do not indicate attacks per se. Although more careful heuristics should be employed in practice to actually feed suspicious events into HEAP, we are able to establish a base line for its validation capabilities nonetheless. We further use HEAP to crosscheck a set of publicly reported real hijacking alarms and demonstrate its practical usefulness in identifying false positives.

### A. Experiment Setup

Our evaluation setup comprises several steps that are repeatedly executed. First, we obtain a full BGP table export holding all prefixes currently present in the global routing system and construct a binary prefix tree such that a tree node holds the date and origin AS of an announcement. To discover emerging subMOAS events, we obtain subsequent BGP messages and update the binary tree accordingly. We consequently extract all *strict* subMOAS events (refer to Subsection II-B4) from the tree that newly appeared in the BGP updates. For these, we apply our filters individually, i.e. we query our graph database for business and resource relations, construct an event-specific AS-level topology, and initiate SSL/TLS measurements for affected hosts that also appear in our ground truth data set. We then retrieve the scan results from successful SSL/TLS handshakes and compare the cryptographic host keys with those from our initial ground truth scan. As outlined earlier, we need to ensure that our scans actually reached a targeted prefix, since our BGP view, respectively the subMOAS events, might be outdated at the time of observation. Thus, we re-evaluate the aforementioned BGP update messages and discard previous scan results for which a subMOAS event changed or vanished during a scan. Note that we accordingly sanitize the data in our ground truth, too, to ensure that no initially scanned SSL/TLS hosts were affected by subMOAS events. This led to the removal of 2,732 hosts.

For the following evaluation, we utilized publicly available BGP data from RouteViews Oregon [26], which provides BGP tables every two hours. As a consequence, we cannot recognize shorter-lived events. This is no inherent limitation: In productive environments, HEAP can be interfaced with a live stream of BGP data, e.g. directly obtained from BGP routers or from services like BGPmon [27].




Fig. 4. subMOAS events observed with our experiment.




Fig. 5. Distribution of subMOAS reoccurrences (CCDF).

TABLE IV  
OVERVIEW OF HEAP RESULTS (COMBINED)

	total	in %
<b>All subMOAS events</b>	<b>14,050</b>	<b>100.0%</b>
IRR analysis	5,699	40.56%
topology reasoning	2,328	16.57%
SSL/TLS scans	2,639	18.78%
<b>Legitimate events (cum.)</b>	<b>7,998</b>	<b>56.93%</b>

### B. Overall Results

Figure 4 shows the frequency of subMOAS events observed during the month of August, 2015. On average, we encountered 88 events every two hours. The minimum number is 1, the maximum number is 388. Figure 5 gives details on subMOAS events that occurred more than once, i.e. concerned the same prefixes and ASes. On average, subMOASes recurred 2.2 times, with a maximum of 169 reoccurrences. In the following, multiple occurrences of identical subMOAS events are considered only once.

During our experiments, we observed a total of 14,050 unique subMOAS events. Our data sources cover 11,222, i.e. 79.87% of these events. Hence, our coverage can still be increased, which suggests that extending HEAP by additional filters can further improve our legitimization results. By feeding the subMOAS events into HEAP, we were able to



TABLE V  
OVERVIEW OF HEAP RESULTS (IRR FILTER)

All subMOAS events: <b>14,050</b>	AfriNIC		APNIC		ARIN		LACNIC		RIPE	
	total	in %	total	in %	total	in %	total	in %	total	in %
<b>Covered subMOAS events</b>	<b>340</b>	<b>2.42%</b>	<b>2,020</b>	<b>14.38%</b>	<b>5,284</b>	<b>37.61%</b>	<b>574</b>	<b>4.09%</b>	<b>3,312</b>	<b>23.57%</b>
valid business relationships	63	0.45%	1,042	7.42%	970	6.90%	n/a	n/a	1,677	11.94%
valid resource holdership	104	0.74%	1,298	9.24%	1,800	12.81%	n/a	n/a	1,971	14.03%
<b>Legitimate events (cum.)</b>	<b>104</b>	<b>0.74%</b>	<b>1,397</b>	<b>9.94%</b>	<b>2,018</b>	<b>14.36%</b>	<b>n/a</b>	<b>n/a</b>	<b>2,452</b>	<b>17.45%</b>

legitimize 56.93%. Table IV presents an overview of individual filter results. Note that an event might be legitimized by multiple filters: in total, we obtain 10,666 legitimate events, which amount to 7,998 distinct cases. At the same time, 5,653 of these cases were legitimized by only a single filter, i.e. each filter contributes unique results. The IRR analysis yields 3,660 unique legitimized cases, followed by SSL/TLS scans with 1,244 cases and our topology reasoning with 749 cases. Overall, we are able to legitimize more than half of all subMOAS events. We will see later on that HEAP performs even better under more realistic conditions, i.e. for alarms relating to networks that are of high value for an attacker.

### C. In-Depth Analysis of the IRR Filter

With the help of the five IRR databases—provided by AfriNIC, APNIC, ARIN, LACNIC, and RIPE—we can legitimize 40.56% of all subMOAS events observed during our analysis period in August, 2015. We identified 5,971 legitimate causes for 11,530 covered events, i.e. for cases where the affected IP prefixes and ASes were registered in one of the IRR databases. Note that some of these resources are registered in multiple databases. Overall, we legitimize a total of 5,699 distinct cases out of 10,500 covered unique events.

Table V shows details on the effectiveness of individual IRR filters at eliminating benign subMOAS events. The highest coverage of events is provided by ARIN (37.61%), while AfriNIC and LACNIC cover less than 5%. At the same time, the ARIN filter legitimates a comparatively low fraction of its covered events due to missing maintainer and RPSL information in the respective IRR data model. In absolute terms, RIPE, ARIN, and APNIC yield the highest number of legitimized events. LACNIC removes all privacy-related information from its IRR database snapshots. As a consequence, none of its covered subMOAS events can be legitimized.

Table II already suggested that filters based on *org*, *import*, and *maintained\_by* relations, i.e. filters utilizing ORGANIZATION, ROUTE and MNTNER objects, show the potential to perform best due to rich relations between these resource objects. Our results confirm this assumption. The most effective filters are based on ORGANIZATION objects in the ARIN database (11.86%), followed by ROUTE objects in the RIPE database (11.17%). Where applicable, maintainer relations are highly effective as well (up to 8.31%). Interestingly, filters that are based on a combination of ORGANIZATION and MNTNER objects contribute least to the overall validation results (0.99% at most). In total, filter rules that aim at identifying business relationships can eliminate 25.17% of all events, while rules that establish confirmation

TABLE VI  
OVERVIEW OF HEAP RESULTS (SSL/TLS FILTER)

SSL/TLS scans	total	in %
	<b>95,486</b>	<b>100.0%</b>
same SSL/TLS key	45,572	47.73%
different SSL/TLS key	13,202	13.83%
no response (port closed)	19,119	20.02%
discarded scan results	17,593	18.42%




Fig. 6. Available SSL/TLS hosts per subMOAS event (CCDF). Events without any such hosts are omitted.

of resource holdership yield 36.39% legitimate events. If we combine them, we find that 40.56% of all subMOAS events (or 54.28% of all covered events) can be legitimized.

### D. In-Depth Analysis of the SSL/TLS Filter

It is worthwhile to study the performance of our SSL/TLS filter in more detail. Table VI shows further information about scans to individual ground truth hosts. In total, we scanned 95,486 SSL/TLS hosts distributed over 3,236 (23.03%) subMOAS events. Note that we discarded 18.42% of the scan results, for which the subMOAS events changed or vanished during the scans. Another 20.02% of our ground truth hosts did not respond to the validation scans. Overall, 47.73% of the retrieved SSL/TLS keys did not change, leading to a total of 2,639 (18.78%) legitimized subMOAS events.

Despite the comparatively high number of unusable scan results, we obtain a relative legitimization rate of 81.55% for covered subMOAS prefixes, i.e. for such prefixes with at least one SSL/TLS-enabled host in our ground truth data set. To elaborate, Figure 6 shows the distribution of available




Fig. 7. Fraction of validated SSL/TLS keys during our experiment.




Fig. 8. Fraction of available SSL/TLS hosts for covered subMOAS events, broken down by port.

SSL/TLS hosts per subMOAS prefix. For 47.13% of all covered events, our ground truth in fact comprises more than three available hosts. 20.65% of these events provide more than ten hosts, and 3.37% of them even more than 100. The maximum number of available hosts is 2,531 with an average of 29.51 hosts per event. These figures actually allow our SSL/TLS filter to be highly robust against outages of individual hosts or services, since it is enough for our technique to confirm that *at least one* cryptographic key on *any* of the affected hosts inside a prefix remains unchanged during a subMOAS event. Figure 7 further indicates that the fractions of unchanged and changing keys shift rather slowly over the time frame of one month. Note that despite the low decline in stable keys, we need to occasionally renew our ground truth data set nonetheless.

Another interesting fact with respect to the legitimization capabilities of our SSL/TLS filter is the set of ports, i.e. network protocols, that contribute to the validation. Figure 8 illustrates that for more than 95% of covered subMOAS events, respectively events with at least one SSL/TLS host available, HTTPS servers can be utilized for the validation. Other protocols like LDAPS, FTPS, XMPPS, and IRCs are apparently ill-suited for our purposes. While adding

TABLE VII  
OVERVIEW OF HEAP RESULTS (COMBINED). EVENTS  
AFFECTING TOP 1 MILLION WEB SITES ONLY

	total	in %
<b>All subMOAS events</b>	<b>849</b>	<b>100.0%</b>
IRR analysis	294	34.63%
topology reasoning	146	17.20%
SSL/TLS scans	576	67.85%
<b>Legitimate events (cum.)</b>	<b>689</b>	<b>81.15%</b>

robustness against outages of HTTPS services for half of the HTTPS-validated events (54.96%), these protocols contribute as few as 75 (2.83%) unique legitimate events. These facts will be taken into account for future re-scans of our ground truth.

### E. A Practical Case Study

So far, we evaluated HEAP with respect to its legitimization capabilities of rather general day-to-day events. To get a more realistic view of its capabilities to identify false positive hijacking alarms in practice, we conduct a case study as follows. We assume that an attacker has little interest in hijacking small and insignificant networks, since the corresponding address space can be easily monitored and, more importantly, has no particular reputation in terms of globally whitelisted IP ranges. Instead, we assume that a real attacker would typically hijack smaller parts of large and popular networks in order to launch and sustain malicious activities. We thus evaluate HEAP with respect to the more well-known networks. To this end, we utilize a list of the top one million web sites provided by *Alexa Inc.* [28]. For each of the domain names in this list, we perform a reverse DNS lookup. Since multiple (sub)domains can be hosted on a single server, we obtain a total of 522,655 distinct IP addresses. We consequently re-assess all subMOAS events during August, 2015 and restrict the input fed into HEAP to those events that affect the aforementioned addresses.

We find that only a small subset of 849 distinct subMOAS cases out of the full set of 14,050 events affect these popular networks. At the same time, we see that the average number of recurring events increases by 13.84%, which indicates intentional use of the subMOAS announcements and supports arguments against misconfiguration or attacks. Hence, one would expect to identify a larger fraction of legitimate subMOAS events. Table VII present the results.

We see that HEAP yields a significantly higher legitimization rate of 81.15% for subMOAS events that relate to the top one million web sites as compared to 56.93% for all observed events (see Table IV). A major reason for this improvement is an increase in coverage of our methodology. The combined filter set now covers 98.82% of the respective events compared to 79.87% in the section before. Most notably, the performance of our SSL/TLS filter is more than three times as high, which is not surprising for much-frequented networks. As a matter of fact, 73.80% of all SSL/TLS hosts in our ground truth data set reside in the popular networks (compare Table VI and Table VIII), but attribute to as few as 6.04% of all subMOAS events. This finding is also reflected by Figure 9, which shows

TABLE VIII  
OVERVIEW OF HEAP RESULTS (SSL/TLS FILTER). EVENTS  
AFFECTING TOP 1 MILLION WEB SITES ONLY

	total	in %
<b>SSL/TLS scans</b>	<b>70,464</b>	<b>100.0%</b>
same SSL/TLS key	31,888	45.25%
different SSL/TLS key	6,508	9.24%
no response (port closed)	17,529	24.88%
discarded scan results	14,539	20.63%




Fig. 9. Available SSL/TLS hosts per subMOAS event (CCDF). Events affecting top 1 million web sites only.

a significant difference in the number of available SSL/TLS hosts per subMOAS event as compared to Figure 6.

The coverage of our IRR filters changes significantly for the *Alexa*-based events (see Table IX). The fraction of subMOAS events that affect the ARIN and APNIC service region increases from 37.61% to 53.71%, and from 14.38% to 25.56% respectively (compare to Table V). The remaining IRR filters, in particular LACNIC and AfriNIC, lose part of their coverage. Altogether, the overall coverage of subMOAS events increases from 74.73% to 93.64%, while the overall legitimization rate slightly decreases from 40.56% to 34.63%.

To put these results into perspective, we use our graph database to identify the responsible registrars for all of the top 1 million web sites, i.e. the respective databases that hold information about corresponding *Alexa* IP prefixes. Most of these prefixes are registered in the ARIN database (99.99%), followed by RIPE (98.86%) and APNIC (90.77%). LACNIC (2.38%) and AfriNIC (0.54%) only account for a small number of these web sites. It is apparent that the largest part of corresponding INETNUM objects is registered in multiple IRR databases. Such networks often relate to several regional subsidiaries of worldwide operating companies under independent administrative control, which possibly explains the slight decrease in our IRR legitimization rate in spite of an increase in coverage.

#### F. Feeding Real Alarms Into HEAP

With HEAP, we intend to provide a framework that enables reliable assessment of arbitrary subprefix hijacking alarms. To demonstrate its effectiveness, we study a set of real alarms

reported by BGPmon.net [16] during August, 2015. This set consists of 85 highly suspicious subprefix hijacking alarms  $\hat{A}$ , each given by

$$\hat{A} = \{vp_v \mid v \in \Sigma_{AS}, p_v \subset \Pi_v\} \cup \{ap'_v \mid a \in \Sigma_{AS}, p'_v \subset p_v\}.$$

During our evaluation, we observed a total of 61 corresponding subMOAS events, for which we applied our filtering scheme. This lower number of observed events compared to the full set of reported alarms results from technical aspects of our experiment design: 7 events lasted for less than two hours, while 9 events were not classified as strict subMOAS (see Subsection V-A). Another 8 events re-occured and were considered only once. For the remaining cases, we retroactively applied our IRR and topology-based filters, while we naturally lack SSL/TLS measurement data from targeted scans.

Table X shows our overall legitimization results. In total, our methodology covered 61 (71.76%) distinct alarms, of which 7 (8.24%) were explicitly identified as false positives. Note that BGPmon.net already provides a highly focused set of alarms, since as few as 85 (0.61%) out of the total number of 14,050 subMOAS events were reported during August, 2015. It is thus highly surprising that these reports still contain nearly 10% false alarms. At the same time, these findings evidence the strength of a cross-validation with HEAP. We plan to provide a public interface to HEAP, which accepts input alarms in the format as specified above supplemented by timestamps of the events. Current and future detection systems may then benefit from our validation scheme as well.

#### G. Summary

With a thorough evaluation of day-to-day anomalies, we established an encouraging base line for practical validation of hijacking alarms: we legitimized 56.93% of these events. In our case study, we further narrowed down the search space for practical hijacking attacks by focusing on networks that host the top one million web sites. Our ability to legitimize 81.15% of corresponding events indicates that our methodology performs even better for such popular networks. These networks may be at higher risk of being attacked due to their good reputation in whitelists. With an analysis of publicly reported hijacking incidents, we demonstrated great practical benefits of our system by identifying nearly 10% of the alarms as false positives. We are thus ready to interface our system with that of fellow researchers to receive and assess their alerts.

Based on our evaluation, we arrived at the following conclusions. First, data obtained from IRR databases, albeit possibly incomplete, is highly useful to assess hijacking alarms in practice. Second, our topology reasoning technique proves to be of equally high effectiveness. Last, but not least, active scans greatly support a reliable assessment of hijacking alarms. The applicability of this approach is remarkably high, which, more importantly, relates to a huge set of SSL/TLS-enabled hosts that remained stable throughout our experiments. We consequently encourage network operators to “opt-in” to HEAP by simply setting up HTTPS servers with unique SSL/TLS keys in their networks—these would be automatically found by our ground truth scans and incorporated into HEAP—ready to be

TABLE IX  
OVERVIEW OF HEAP RESULTS (IRR FILTER). EVENTS AFFECTING TOP 1 MILLION WEB SITES ONLY

All subMOAS events: <b>849</b>	AfriNIC		APNIC		ARIN		LACNIC		RIPE	
	total	in %	total	in %	total	in %	total	in %	total	in %
<b>Covered subMOAS events</b>	<b>6</b>	<b>0.71%</b>	<b>217</b>	<b>25.56%</b>	<b>456</b>	<b>53.71%</b>	<b>22</b>	<b>2.59%</b>	<b>139</b>	<b>16.37%</b>
valid business relationships	4	0.47%	62	7.30%	16	1.89%	0	0.00%	57	6.74%
valid resource holdership	4	0.47%	121	14.25%	59	7.00%	0	0.00%	73	8.60%
<b>Legitimate events (cum.)</b>	<b>4</b>	<b>0.47%</b>	<b>138</b>	<b>16.25%</b>	<b>62</b>	<b>7.30%</b>	<b>0</b>	<b>0.00%</b>	<b>98</b>	<b>11.54%</b>

TABLE X  
HEAP CROSS-CHECK OF BGPmon.net HIJACKING ALARMS [16]

reported alarms: <b>85</b>	IRR analysis		Topology reasoning		SSL/TLS scans		total (cum.)	
	total	in %	total	in %	total	in %	total	in %
covered alarms	60	70.59%	3	3.53%	1	1.18%	61	71.76%
<b>false positives</b>	<b>6</b>	<b>7.06%</b>	<b>1</b>	<b>1.18%</b>	<b>0</b>	<b>0.00%</b>	<b>7</b>	<b>8.24%</b>

used for validation scans in case of an alarm. In our evaluation, we further observed striking differences in the deployment of SSL/TLS, which led to improvements to our plans for regular ground truth scans in the future.

## VI. CONCLUSION AND OUTLOOK

In this paper, we introduced a novel approach to formalize characteristics of Internet routing. Applied to BGP hijacking, it is suitable to precisely formulate, classify, and evaluate different kinds of attacks. We utilized this model to assess impact and traceability of subprefix hijacking attacks. A concept of general nature, it may serve to establish a basis to address future research questions on Internet routing.

Based on our formal attacker model, we derived HEAP, an extensible filtering system that combines several data sources in order to reliably assess the validity of hijacking alarms. An automated reasoning technique for given routing anomalies, it lends itself well to integration with the state-of-the-art and future hijacking detection systems, in particular to crosscheck and narrow down their number of false alarms. In our evaluation, we thoroughly analyzed the applicability of our approach and demonstrated its usefulness in practice by revealing a significant number of false positives in a set of well-established hijacking reports. We intend to grow our framework into a public service that makes its data available on a continuous basis. We invite researchers to feed our system with their conjectural alerts and to further extend the system by resourceful data sources.

## ACKNOWLEDGEMENTS

The authors thank the reviewers for their constructive comments.

## REFERENCES

- [1] H. Ballani, P. Francis, and X. Zhang, "A study of prefix hijacking and interception in the Internet," in *Proc. ACM SIGCOMM Int. Conf.*, 2007, pp. 265–276.
- [2] J. Schlamp, G. Carle, and E. W. Biersack, "A forensic case study on as hijacking: The attacker's perspective," *ACM SIGCOMM Comput. Commun. Rev.*, vol. 43, no. 2, pp. 5–12, Apr. 2013.
- [3] S. Kent, C. Lynn, and K. Seo, "Secure border gateway protocol (S-BGP)," *IEEE J. Sel. Areas Commun.*, vol. 18, no. 4, pp. 582–592, Apr. 2000.
- [4] G. Huston and R. Bush, "Securing BGP and SIDR," *IETF J.*, vol. 7, no. 1, pp. 8–13, 2011.
- [5] M. Lad, D. Massey, D. Pei, Y. Wu, B. Zhang, and L. Zhang, "PHAS: A prefix hijack alert system," in *Proc. 15th USENIX Secur. Symp. (USENIX-SS)*, 2006, pp. 153–166.
- [6] J. Qiu, L. Gao, S. Ranjan, and A. Nucci, "Detecting bogus BGP route information: Going beyond prefix hijacking," in *Proc. 3rd Int. Conf. Secur. Privacy Commun. Netw. (SecureComm)*, Sep. 2007, pp. 381–390.
- [7] C. Zheng, L. Ji, D. Pei, J. Wang, and P. Francis, "A light-weight distributed scheme for detecting IP prefix hijacks in real-time," in *Proc. ACM SIGCOMM Int. Conf.*, 2007, pp. 277–288.
- [8] Z. Zhang, Y. Zhang, Y. C. Hu, Z. M. Mao, and R. Bush, "iSPY: Detecting IP prefix hijacking on my own," in *Proc. ACM SIGCOMM Int. Conf.*, 2008, pp. 327–338.
- [9] X. Hu and Z. M. Mao, "Accurate real-time identification of IP prefix hijacking," in *Proc. IEEE Symp. Secur. Privacy (IEEESSP)*, May 2007, pp. 3–17.
- [10] X. Shi, Y. Xiang, Z. Wang, X. Yin, and J. Wu, "Detecting prefix hijackings in the Internet with argus," in *Proc. ACM SIGCOMM Internet Meas. Conf. (IMC)*, 2012, pp. 15–28.
- [11] J. Schlamp et al., "Investigating the nature of routing anomalies: Closing in on subprefix hijacking attacks," in *Proc. Int. Workshop Traffic Monitor. Anal. (TMA)*, Apr. 2015, pp. 173–187.
- [12] Q. Jacquemart, G. Urvoy-Keller, and E. Biersack, "A longitudinal study of BGP MOAS prefixes," in *Proc. Int. Workshop Traffic Monitor. Anal. (TMA)*, Apr. 2014, pp. 127–138.
- [13] J. Schlamp, J. Gustafsson, M. Wählisch, T. C. Schmidt, and G. Carle, "The abandoned side of the Internet: Hijacking Internet resources when domain names expire," in *Proc. Int. Workshop Traffic Monitor. Anal. (TMA)*, Apr. 2015, pp. 188–201.
- [14] T. Qiu, L. Ji, D. Pei, J. Wang, J. Xu, and H. Ballani, "Locating prefix hijackers using LOCK," in *Proc. USENIX Secur. Symp. (USENIX-SS)*, 2009, pp. 135–150.
- [15] J. Li, T. Ehrenkranz, and P. Elliott, "Buddyguard: A buddy system for fast and reliable detection of IP prefix anomalies," in *Proc. 20th IEEE Int. Conf. Netw. Protocols (ICNP)*, Oct./Nov. 2012, pp. 1–10.
- [16] BGPmon Network Solutions Inc. (2015). *Routing Anomalies*. [Online]. Available: <http://www.bgstream.com>
- [17] M. Tahara, N. Tateishi, T. Oimatsu, and S. Majima, "A method to detect prefix hijacking by using ping tests," in *Proc. 11th Asia-Pacific Netw. Oper. Manage. Symp. (APNOMS)*, 2008, pp. 390–398.
- [18] J. W. Mickens, J. R. Douceur, W. J. Bolosky, and B. D. Noble, "StrobeLight: Lightweight availability mapping and anomaly detection," in *Proc. USENIX Annu. Tech. Conf. (USENIX-ATC)*, 2009, pp. 1–14.
- [19] S.-C. Hong, H.-T. Ju, and J. W. Hong, "IP prefix hijacking detection using idle scan," in *Proc. 12th Asia-Pacific Netw. Oper. Manage. Symp. (APNOMS)*, 2009, pp. 395–404.
- [20] S.-C. Hong, J. W.-K. Hong, and H. Ju, "IP prefix hijacking detection using the collection of AS characteristics," in *Proc. 13th Asia-Pacific Netw. Oper. Manage. Symp. (APNOMS)*, 2011, pp. 1–7.
- [21] P.-A. Vervier and O. Thonnard, "SpamTracer: How stealthy are spammers?" in *Proc. Int. Workshop Traffic Monitor. Anal. (TMA)*, 2013, pp. 3477–3482.
- [22] A. Khan, H.-C. Kim, T. Kwon, and Y. Choi, "A comparative study on IP prefixes and their origin as in BGP and the IRR," *ACM SIGCOMM Comput. Commun. Rev.*, vol. 43, no. 3, pp. 16–24, Jul. 2013.
- [23] R. Holz, J. Amann, O. Mehani, M. Wachs, and M. A. Kaafar, "TLS in the wild—An Internet-wide analysis of TLS-based protocols for electronic communication," in *Proc. ISOC Symp. Netw. Distrib. Syst. Secur. (NDSS)*, 2016.



- [24] A. Pilosov and T. Kapela, "Stealing the Internet: An Internet-scale man in the middle attack," in *Proc. DEFCON*, 2008.
- [25] C. Hepner and E. Zmijewski, "Defending against BGP man-in-the-middle attacks," in *Proc. BlackHat*, 2009.
- [26] D. Meyer. (2005). *University of Oregon RouteViews Project*. [Online]. Available: <http://www.routeviews.org>
- [27] H. Yan, R. Oliveira, K. Burnett, D. Matthews, L. Zhang, and D. Massey, "BGPmon: A real-time, scalable, extensible monitoring system," in *Proc. Cybersecur. Appl. Technol. Conf. Homeland Secur. (CATCH)*, 2009, pp. 212–223.
- [28] Alexa Internet, Inc. (2015). *The Top 1 Million Sites on the Web*. [Online]. Available: <http://s3.amazonaws.com/alexa-static/top-1m.csv.zip>




**Johann Schlamp** received the degree in computer science with a minor in mathematics from the Technical University of Munich in 2009. In his Ph.D. thesis, he conducted an evaluation of architectural threats to Internet routing. Since then, he has been with the Chair of Network Architectures and Services. He is currently a Research Associate with the Technical University of Munich, Germany. His research interests lie in interdomain routing, Internet protocols, and network measurements.



IETF and is an author of two RFCs.

**Ralph Holz** received the Ph.D. degree from the Technical University of Munich, Germany. He led the research that first analysed and demonstrated the security shortcomings of X.509 for the Web with the Technical University of Munich. He is currently a Lecturer in networks and security with The University of Sydney and a Conjoint Researcher with Data61|CSIRO, Australia's premier innovation hub. His current research agenda is empirical security, in particular global-scale analyses of Internet service deployments and technology. He contributes to the



**Quentin Jacquemart** received the master's degree in computer science from the University of Liège, Belgium, in 2011, and the Ph.D. degree from Telecom ParisTech, France, in 2015. From 2011 to 2015, he was part of the Networking and Security Group with Eurécom, Sophia Antipolis, France. He is currently a Post-Doctoral Research Fellow with the SigNet Group, I3S Laboratory, University of Nice Sophia Antipolis, France. His main research interest is the security of computer networks and of their protocols.



the Department of Informatics, Technical University of Munich, Germany, holding the Chair of Network Architectures and Services.

**Georg Carle** received the degree in electrical engineering from the University of Stuttgart, Stuttgart, Germany, the M.Sc. degree from Brunel University, London, U.K., and the Ph.D. degree from the University of Karlsruhe. He had a stay with Telecom ParisTech, Paris, France. He was a Post-Doctoral Scientist with Eurécom, Sophia Antipolis, France, and the Fraunhofer Institute for Open Communication Systems, Berlin, Germany, prior to being appointed as a Professor with the University of Tübingen, Germany. He is currently a Professor with



**Ernst W. Biersack** received the degree in computer science and the Ph.D. degree from the Technical University of Munich, Germany. He had a stay at University of North Carolina, Chapel Hill, NC, USA. He received the Habilitation à Diriger des Recherches degree from the University of Nice, France. From 1989 to 1992, he was a member of the Technical Staff with the Computer Communications Research Group, Bell Communications Research, Morristown, NJ, USA. From 1992 to 2014, he was a Professor with Eurécom, Sophia Antipolis, France.